

Data Mining :: Unit-3

(Classification – Nearest Neighbor Classifier)

Er. Dinesh Baniya Kshatri
(Lecturer)

Department of Electronics and Computer Engineering
Institute of Engineering, Thapathali Campus

Nearest Neighbor Classifier

- Tell me who your friends are and I'll tell you who you are
- A **new example** is assigned to the most common class among the (K) examples that are most similar to it.

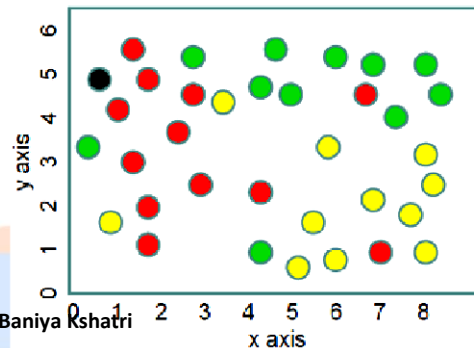


Prepared by: Er. Dinesh Baniya Kshatri

2

Example: 1-Nearest Neighbor Classifier [1]

- Suppose we have a problem where:
 - We have three classes (red, green, yellow).
 - Each pattern is a two-dimensional vector.
- Suppose that the training data is given below:
- Suppose we have a test pattern v , shown in black.
- How is v classified by the nearest neighbor classifier?

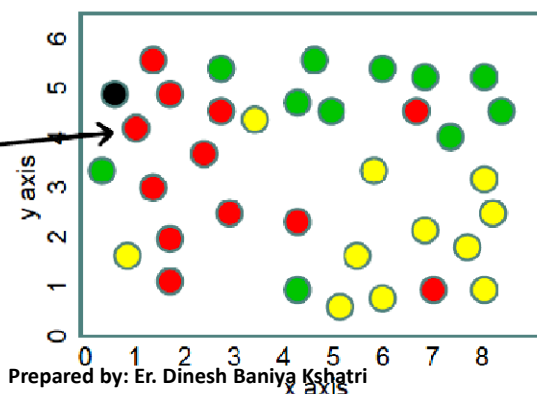


Prepared by: Er. Dinesh Baniya Kshatri

3

Example: 1-Nearest Neighbor Classifier [2]

- Suppose we have a problem where:
 - We have three classes (red, green, yellow).
 - Each pattern is a two-dimensional vector.
- Suppose that the training data is given below:
- Suppose we have a test pattern v , shown in black.
- How is v classified by the nearest neighbor classifier?
- $C(v)$:
- Class of $C(v)$: red.
- Therefore, v is classified as red.

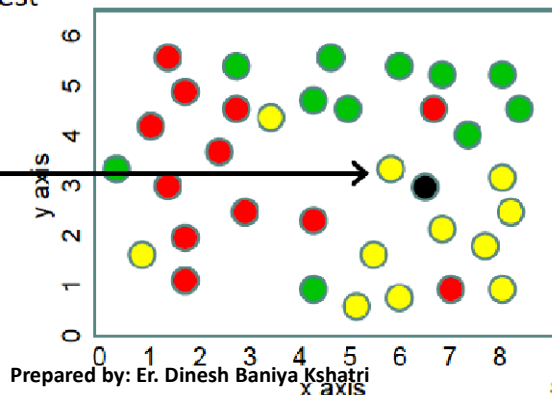


Prepared by: Er. Dinesh Baniya Kshatri

4

Example: 1-Nearest Neighbor Classifier [3]

- Suppose we have a problem where:
 - We have three classes (red, green, yellow).
 - Each pattern is a two-dimensional vector.
- Suppose that the training data is given below:
- Suppose we have another test pattern v , shown in black.
- How is v classified by the nearest neighbor classifier?
- $C(v)$: —————→
- Class of $C(v)$: yellow.
- Therefore, v is classified as yellow.




Prepared by: Er. Dinesh Baniya Kshatri

5

Example: 3-Nearest Neighbor Classifier

- Instead of classifying the test pattern based on its nearest neighbor, we can take more neighbors into account.
- This is called k -nearest neighbor classification.
- Using more neighbors can help avoid mistakes due to noisy data.
- In the example shown on the figure, the test pattern is in a mostly "yellow" area, but its nearest neighbor is red.
- If we use the 3 nearest neighbors, 2 of them are yellow.



Prepared by: Er. Dinesh Baniya Kshatri

6

त्रिभुवन विश्वविद्यालय The kNN Algorithm (Pseudo-code)

- Determine parameter K
- Calculate the distance between the test instance and all the training instances
- Sort the distances and determine K nearest neighbors
- Gather the labels of the K nearest neighbors
- Use simple majority voting or weighted voting.

Prepared by: Er. Dinesh Baniya Kshatri

7

त्रिभुवन विश्वविद्यालय The kNN Algorithm (Mathematical Formalism)

- Let F be a distance function defined in \mathbb{X} .
 - F assigns a distance to every pair v_1, v_2 of objects in \mathbb{X} .
 - Let x_1, x_2, \dots, x_N be training examples.
 - The nearest neighbor classifier classifies any pattern v as follows:
 - Find the training example $C(v)$ that is the nearest neighbor of x (has the shortest distance to v among all training data).
- $$C(v) = \operatorname{argmin}_{x \in \{x_1, \dots, x_N\}} (F(v, x))$$
- Return the class label of $C(v)$.
 - In short, each test pattern v is assigned the class of its nearest neighbor in the training data.

• Let, \mathbb{X} be the space of all possible patterns for some classification problem

Prepared by: Er. Dinesh Baniya Kshatri

8

Properties of the kNN Algorithm

- This algorithm belongs to the class of “lazy” algorithms. There is no process of learning or training. The examples are simply stored as the data is collected.
- The difficulty comes at classification stage. We need to calculate n distances and find best K data points.
- How to choose K
 - too small: the method might be inaccurate and sensitive to noise
 - too large: the method is more robust but may lose sensitivity to changes in the feature space
 - solution? Try a few K 's and find optimum based on the estimated accuracy of the predictor. There exist formal algorithms to select K .

Prepared by: Er. Dinesh Baniya Kshatri

9

kNN: How to Choose k ? – [1]


- Rule of thumb is $k = \text{sqrt}(n)$, n is number of examples
- In practice, $k = 1$ is often used for efficiency, but can be sensitive to “noise”
- larger k may improve performance, but too large k destroys *locality*, i.e. end up looking at samples that are not neighbors
- cross-validation (study later) may be used to choose k

Prepared by: Er. Dinesh Baniya Kshatri

10


kNN: How to Choose k? – [2]

1 NN



every example in the blue shaded area will be misclassified as the **blue** class

3 NN



every example in the blue shaded area will be classified correctly as the **red** class


Prepared by: Er. Dinesh Baniya Kshatri

11

Nearest Neighbor Classifiers (Classification Requirements)

Unknown record

$k=3$




Requires **three** inputs:

1. The set of stored training samples
2. Distance metric to compute distance between samples
3. The value of k , i.e., the number of nearest neighbors to retrieve

Prepared by: Er. Dinesh Baniya Kshatri

12

Nearest Neighbor Classifiers (Classification Process)




To classify **unknown** sample:

1. Compute distance to other training records
2. Identify k nearest neighbors
3. Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)


Prepared by: Er. Dinesh Baniya Kshatri

13


Definition of Nearest Neighbor



(a) 1-nearest neighbor



(b) 2-nearest neighbor




(c) 3-nearest neighbor

k -nearest neighbors of a sample x are datapoints that have the k smallest distances to x .

Prepared by: Er. Dinesh Baniya Kshatri

14

Concept of Voronoi Diagrams – [1]



15

Concept of Voronoi Diagrams – [2]

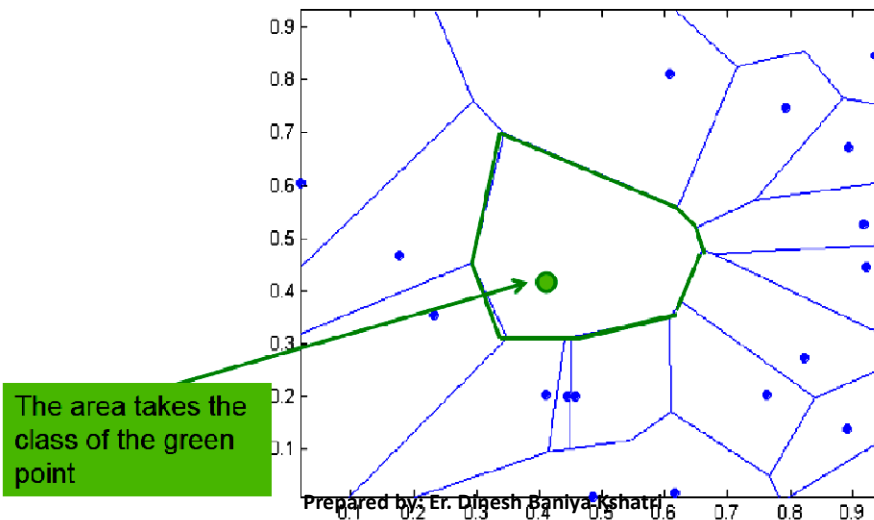
- **Property 1:**
 - A Voronoi diagram divides a space into disjoint polygons where the nearest neighbor of any point inside a polygon is the generator of the polygon.
- **Property 2:**
 - Each Voronoi edge is a segment of the perpendicular bisector of a pair of generators.
- **Property 3:**
 - Each Voronoi edge is shared by two Voronoi polygons and average number of Voronoi edges per Voronoi polygon is at most 6. This means that each generator has 6 adjacent generators at most.

Prepared by: Er. Dinesh Baniya Kshatri

16

Concept of Voronoi Diagrams – [3]

Voronoi Diagram defines the classification boundary



17

Distance between Neighbors

- Each example is represented with a set of numerical attributes



John:
Age=35
Income=95
No. of credit cards=3



Rachel:
Age=41
Income=215
No. of credit cards=2

- "Closeness" is defined in terms of the Euclidean distance between two examples.
 - The Euclidean distance between $X=(x_1, x_2, x_3, \dots, x_n)$ and $Y=(y_1, y_2, y_3, \dots, y_n)$ is defined as:







$$D(X, Y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Distance (John, Rachel) = $\sqrt{(35-41)^2 + (95-215)^2 + (3-2)^2}$

18

Class Work – Question

(Use kNN, with $k = 3$ to find David's Response)







Customer	Age	Income	No. credit cards	Response
John 	35	35	3	No
Rachel 	22	50	2	Yes
Hannah 	63	200	1	No
Tom 	59	170	1	No
Nellie 	25	40	4	Yes
David 	37	50	2	?

Prepared by: Er. Dinesh Baniya Kshatri

19

Class Work – Solution

(Distances from David & David's Response)

Customer	Age	Income	No. cards	Response	Distance from David
John 	35	35	3	No	$\text{sqrt} [(35-37)^2 + (35-50)^2 + (3-2)^2] = 15.16$
Rachel 	22	50	2	Yes	$\text{sqrt} [(22-37)^2 + (50-50)^2 + (2-2)^2] = 15$
Hannah 	63	200	1	No	$\text{sqrt} [(63-37)^2 + (200-50)^2 + (1-2)^2] = 152.23$
Tom 	59	170	1	No	$\text{sqrt} [(59-37)^2 + (170-50)^2 + (1-2)^2] = 122$
Nellie 	25	40	4	Yes	$\text{sqrt} [(25-37)^2 + (40-50)^2 + (4-2)^2] = 15.74$
David 	37	50	2	Yes	

Prepared by: Er. Dinesh Baniya Kshatri

20

Distance between Instances (Numeric Features)

- Euclidean distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_2 = \sqrt{\sum_{i=1}^n (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

- Manhattan distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_1 = \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|$$

- L_p -norm

- Euclidean = L_2

- Manhattan = L_1

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_p = \left(\sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|^p \right)^{\frac{1}{p}}$$

Prepared by: Er. Dinesh Baniya Kshatri

21

Distance between Instances (Symbolic / Categorical Features)

Most common distance is the *Hamming distance*

- Number of bits that are different
- Or: Number of features that have a different value
- Also called the *overlap*
- Example:

\mathbf{x}_1 : {Shape=Triangle, Color=Red, Location=Left, Orientation=Up}

\mathbf{x}_2 : {Shape=Triangle, Color=Blue, Location=Left, Orientation=Down}

Hamming distance = 2

Prepared by: Er. Dinesh Baniya Kshatri

22

kNN: Feature Weighting – [1]

- So far we assumed we use Euclidian Distance to find the nearest neighbor
- Euclidean distance treats each feature as equally important
- However some features (dimensions) may be much more discriminative than other features

Prepared by: Er. Dinesh Baniya Kshatri

23

kNN: Feature Weighting – [2]

- Scale each feature by its importance for classification

$$D(a, b) = \sqrt{\sum_k w_k (a_k - b_k)^2}$$

- Can use our prior knowledge about which features are more important
- Can learn the weights w_k using cross-validation (to be covered later)

Prepared by: Er. Dinesh Baniya Kshatri

24

Need for Normalization of Variables



John:
Age=35
Income=95K
No. of credit cards=3



Rachel:
Age=41
Income=215K
No. of credit cards=2

Distance (John, Rachel) = $\sqrt{[(35-41)^2 + (95,000-215,000)^2 + (3-2)^2]}$

- Distance between neighbors could be dominated by some attributes with relatively large numbers (e.g., income in our example). Important to normalize some features (e.g., map numbers to numbers between 0-1)

Example: Income

Highest income = 500K

Davis's income is normalized to 95/500, Rachel income is normalized to 215/500, etc.)

Prepared by: Er. Dinesh Baniya Kshatri

25

kNN: Feature Normalization – [1]

- First feature takes values between 1 to 2
- Second feature takes values between 100 to 200
- Idea:** normalize features to be on the same scale
- Different normalization approaches
- Linearly scale the range of each feature to be, say, in range [0,1]

$$f_{new} = \frac{f_{old} - f_{old}^{min}}{f_{old}^{max} - f_{old}^{min}}$$

Prepared by: Er. Dinesh Baniya Kshatri

26

kNN: Feature Normalization – [2]

- Linearly scale to **0** mean variance **1**:
- If Z is a random variable of mean m and variance σ^2 , then $(Z - m)/\sigma$ has mean **0** and variance **1**
- For each feature f let the new rescaled feature be

$$f_{new} = \frac{f_{old} - \mu}{\sigma}$$

Prepared by: Er. Dinesh Baniya Kshatri

27

Example: Normalizing Variables – [1]

- Suppose that your test patterns are 2-dimensional vectors, representing stars.
 - The first dimension is surface temperature, measured in Fahrenheit.
 - Your second dimension is mass, measured in pounds.
- The surface temperature can vary from 6,000 degrees to 100,000 degrees.
- The mass can vary from 10^{29} to 10^{32} .
- Does it make sense to use the Euclidean distance or the Manhattan distance here?

Prepared by: Er. Dinesh Baniya Kshatri

28

Example: Normalizing Variables – [2]

- Does it make sense to use the Euclidean distance or the Manhattan distance in example of previous slide ?
- No. These distances treat both dimensions equally, and assume that they are both measured in the same units.
- Applied to these data, the distances would be dominated by differences in mass, and would mostly ignore information from surface temperatures.

Prepared by: Er. Dinesh Baniya Kshatri

29

Example: Normalizing Variables – [3]

- It would make sense to use the Euclidean or Manhattan distance, if we first normalized dimensions, so that they contribute equally to the distance.
- How can we do such normalizations?
- There are various approaches. Two common approaches are:
 - Translate and scale each dimension so that its minimum value is 0 and its maximum value is 1.
 - Translate and scale each dimension so that its mean value is 0 and its standard deviation is 1.

Prepared by: Er. Dinesh Baniya Kshatri

30

Example: Normalizing Variables – [4]

Original Data			Normalized Data: Min = 0, Max = 1		Normalized Data: Mean = 0, std = 1	
Object ID	Temp. (F)	Mass (lb.)	Temp.	Mass	Temp.	Mass
1	4700	1.5×10^{30}	0.0000	0.0108	-0.9802	-0.6029
2	11000	3.5×10^{30}	0.1525	0.0377	-0.5375	-0.5322
3	46000	7.5×10^{31}	1.0000	1.0000	1.9218	1.9931
4	12000	5.0×10^{31}	0.1768	0.6635	-0.4673	1.1101
5	20000	7.0×10^{29}	0.3705	0.0000	0.0949	-0.6311
6	13000	2.0×10^{30}	0.2010	0.0175	-0.3970	-0.5852
7	8500	8.5×10^{29}	0.0920	0.0020	-0.7132	-0.6258
8	34000	1.5×10^{31}	0.7094	0.1925	1.0786	-0.1260

Prepared by: Er. Dinesh Baniya Kshatri

31

Nearest Neighbor Search

- The problem of finding the nearest neighbors of a pattern is called “nearest neighbor search”.
- Suppose that we have N training examples.
- Suppose that each example is a D -dimensional vector.
- What is the time complexity of finding the nearest neighbors of a test pattern?
- $O(ND)$.
 - We need to consider each dimension of each training example.

Prepared by: Er. Dinesh Baniya Kshatri

32

kNN: Computational Complexity

- Basic kNN algorithm stores all examples
- Suppose we have n examples each of dimension d
- $O(d)$ to compute distance to one example
- $O(nd)$ to find one nearest neighbor
- $O(knd)$ to find k closest examples
- Thus total complexity is $O(knd)$
- Very expensive for a large number of samples
- But we need a large number of samples for kNN to work well!

Prepared by: Er. Dinesh Baniya Kshatri

33

Nearest Neighbor Search (Indexing Methods)

- As we just mentioned, measuring the distance between the test pattern and each training example takes $O(ND)$ time.
- This method of finding nearest neighbors is called “brute-force search”, because we go through all the training data.
- There are methods for finding nearest neighbors that are sublinear to N (even logarithmic, at times), but exponential to D .
- Can you think of an example?
 - Binary search (applicable when $D = 1$) takes $O(\log N)$ time.

Prepared by: Er. Dinesh Baniya Kshatri

34

Nearest Neighbor Search (Indexing Methods)

- In some cases, faster algorithms exist, however, are approximate.
 - They do not guarantee finding the true nearest neighbor all the time.
 - They guarantee that they find the true nearest neighbor with a certain probability.

Prepared by: Er. Dinesh Baniya Kshatri

35

kNN: How Well does it Work?

- kNN is simple and intuitive, but does it work?
- Theoretically, the best error rate is the Bayes rate E^*
 - Bayes error rate is the best (smallest) error rate a classifier can have, for a given problem, but we do not study it in this course
- Assume we have an unlimited number of samples
- kNN leads to an error rate greater than E^*
- But even for $k=1$, as $n \rightarrow \infty$, it can be shown that kNN error rate is smaller than $2E^*$
- As we increase k , the upper bound on the error gets better, that is the error rate (as $n \rightarrow \infty$) for the **kNN** rule is smaller than cE^* , with smaller c for larger k
- **If we have lots of samples, kNN works well**

Prepared by: Er. Dinesh Baniya Kshatri

36